EE239AS Project 5

Popularity Prediction on Twitter

Winter 2017

i) Popularity Prediction

Introduction: A useful practice in social network analysis is to predict future popularity of a subject or event. Twitter, with its public discussion model, is a good platform to perform such analysis. With Twitter's topic structure in mind, the problem can be stated as: knowing current (and previous) tweet activity for a hashtag, can we predict its tweet activity in the future? More specifically, can we predict if it will become more popular and if so by how much? In this project, we will try to formulate and solve an instance of such problems.

The available Twitter data is collected by querying popular hashtags related to the 2015 Super Bowl spanning a period starting from 2 weeks before the game to a week after the game. We will use data from some of the related hashtags to train a regression model and then use the model to making predictions for other hashtags. To train the model, you need to prepare training sets out of the data, extract selected features for them, and then fit the regression model on it. The regression model will try to fit a curve through observed values of features and outcomes to create a predictor for new samples. Designing and choosing good features is one of the most important steps in this process and is essential to getting a more accurate system. There are examples of such analysis and useful features in literature[1](You should look into the literature for this). You will be given training data to create the model, and test data to make predictions. The test data consists of tweets containing a hashtag in a specified time window, and you will use your model to predict number of tweets containing the hashtag posted within one hour immediately following the given time window.

1) Download the training tweet data[2] and calculate these statistics for each hashtag: average number of tweets per hour, average number of followers of users posting the tweets, and average number of retweets. Plot "number of tweets in hour" over time for #SuperBowl and #NFL (a histogram with 1-hour bins).

The tweets are stored in separate files for different hashtags and files are named as `tweet_[#hashtag].txt`.The tweet file contains one tweet in each line and tweets are sorted with respect to their posting time. Each tweet is a JSON string that you can load in Python as a dictionary.

---

[1] http://arxiv.org/abs/1401.2018
[2] https://ucla.box.com/s/nv9td9kvvfvg3tya0dlvbs1kn5o87gmv

2) Fit a Linear Regression model using 5 features to predict number of tweets in the next hour, with features extracted from tweet data in the previous hour. The features you should use are: number of tweets, total number of retweets, sum of the number of followers of the users posting the hashtag, maximum number of followers of the users posting the hashtag, and time of the day (which could take 24 values that represent hours of the day with respect to a given time reference). Explain your model's training accuracy and the significance of each feature using the t-test and P-value results of fitting the model.

Hint: You can create time windows from the data to extract features. Each window will provide a sample for your regression model. E.g. You can divide the data in 1-hour windows and use features from each 1-hour (or n-hour) window to predict number of tweets for the next 1-hour window.

3) Design a regression model using any features from the papers you find or other new features you may find useful for this problem. Fit your model on the data and report fitting accuracy and significance of variables. For the top 3 features in your measurements, draw a scatter plot of predictant (number of tweets for next hour) versus feature value, using all the samples you have extracted and analyze it.

4) Split the feature data (your set of *(features,predictant)* pairs for windows) into 10 parts to perform cross-validation. Run 10 tests, each time fitting your model on 9 parts and predicting the number of tweets for the 1 remaining part. Calculate the average prediction error $\left| N_{predicted} - N_{real} \right|$ over samples in the remaining part, and then average these values over the 10 tests.

Since we know the Super Bowl's date and time, we can create different regression models for different periods of time. First, when the hashtags haven't become very active, second, their active period, and third, after they pass their high-activity time. Train 3 regression models for these time periods (The times are all in PST):

1. Before Feb. 1, 8:00 a.m.

2. Between Feb. 1, 8:00 a.m. and 8:00 p.m.

3. After Feb. 1, 8:00 p.m.

Report cross-validation errors for the 3 different models. Note that you should do the 90-10% splitting for each model within its specific time window. I.e. Only use data within one of the 3 periods for training and testing each time, so for each period you will run 10 tests.

5) Download the test data[3] and run your model to make predictions for the next hour in each case. Each file in the test data contains a hashtag's tweets for a 6-hour window. The file name shows sample

---

[3] https://ucla.box.com/s/ojvvthudugp9d2gze5nuep9ogwjydnur

number followed by the period number the data is from. E.g. a file named `sample5_period2.txt` contains tweets for a 6-hour window that lies in the 2nd time period described in part 4.

Report your predicted number of tweets for the next hour of each sample window.

ii) Fan Base Prediction

6) Often the textual content of a tweet reveals some information about the author. It is even more revealing if the users tweeting on the same topic have different or even opposing views about it. As such, the tweets posted by fans of different teams during a sport game describe similar events in different terms and sentiments. Recognizing that supporting a sport team has a lot to do with the user location, we try to use the textual content of the tweet posted by a user to predict her location. In order to make the problem more specific, let us consider all the tweets including #superbowl, posted by the users whose specified location is either in the sate of Washington or Massachusetts. For example, in order to include all the tweets with the author located in the state of Washington, we consider the tweets that include the following substrings in the location field:

```
Seattle, Washington
Washington
WA
Seattle, WA
Kirkland, Washington
…
```

Train a classifier to predict the location of the author of a tweet given only the textual content of the tweet using the techniques you learnt in project 2. Try different classification algorithms and in your submission, plot the ROC curve, report the confusion matrix and calculate the accuracy, recall and precision of your classifier.

iii) Define your own Project

7) The dataset in hands is very rich as there is a lot of metadata to a tweet. Be creative and propose a new problem (something interesting that can be inferred from this dataset) other than popularity prediction. You can look into the literature of Twitter data analysis to get some ideas.

Go ahead and implement your idea and show that it works. Even if you cannot implement the task fully, you get credit for the novelty and partial implementation of it.

Submission: This project is due on Wednesday March 22 at 11:59pm. Late submissions are subject to 10% penalty each day past the deadline. Please submit a zip file containing your report, and your

codes with a readme file on how to run your code to ee219.winter2017@gmail.com. The zip file should be named as "Project4_UID1_UID2_..._UIDn.zip" where UIDx are student ID numbers of the team members. If your zip file is too large to email you can send it to the dropbox account associated with this email. You can use any programming language but there will be samples of Python code given to you for crawling and parsing that you can use and modify. If you had any questions you can send an email to the same address.